

CAAVO HackerEarth Computer Vision Challenge

Problem Statement: To classify images of the Garment type

APPROACH

- Download the dataset (train, test, sample_submission.csv)
- Explore the dataset through some visualizations to make sure what we are dealing with. As the class labels are not given, they are given as 0 to 14. It makes a lot of sense to go through the images in each folder (0 to 14) to get an idea of what we are dealing with.
- Split the Downloaded train into validation set and store them in a separate valid folder with each class with its class label as a folder under the valid folder.
- Use a pre-trained network of Resnet34 that is trained using Imagenet images.
- I am using Fastai wrapper of a PyTorch Deep Learning framework created by Jamie Howard.
- Build a baseline model with Fully connected layer training and benchmark the score in the leaderboard
- Introduce Image augmentation for over-fitting problems
- Try Unfreezing the previous layers of Resnet34 architecture and train them with differential learning rates for layers
- Experiment changing the size of the images fed, batch size and learning rates

Feature Engineering

As it an image classification problem, nothing much of a feature engineering required here but image augmentation.

I tried Image augmentation with varying zoom parameters and other stuff.

I also split the train data into valid data for feeding into the resnet34 architecture

TOOLS

PyTorch Deep Learning Framework

Fastai wrapper

Pandas

Scikit-learn

Python

Jupyter for Python notebooks

Paperspace Cloud GPU for processing

SOURCE CODE

Importing the libraries

```
%reload_ext autoreload
%autoreload 2
%matplotlib inline
```

```
from fastai.imports import *
from fastai.transforms import *
from fastai.conv_learner import *
from fastai.model import *
from fastai.dataset import *
from fastai.sgdr import *
from fastai.plots import *
```

Specifying the data to be used

```
PATH = "/home/paperspace/projects/caavo/with_valid_dataset/"
sz=64
```

Specifying the architecture

```
arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch,
sz),test_name='test')
```

Specifying the learner

```
learn = ConvLearner.pretrained(arch, data, precompute=True)
```

Learning the model

```
learn.fit(0.01, 2)
```

Image Augmentation

```
tfms = tfms_from_model(resnet34, sz, aug_tfms=transforms_side_on,
max_zoom=1.1)
```

```
def get_augs():
```

```
    data = ImageClassifierData.from_paths(PATH, bs=2, tfms=tfms, num_workers=1)
    x,_ = next(iter(data.aug_dl))
    return data.trn_ds.denorm(x)[1]
```

Loading the data and the learner

```
data = ImageClassifierData.from_paths(PATH, tfms=tfms, test_name='test')
learn = ConvLearner.pretrained(arch, data, precompute=True)
```

Fitting the model

```
learn.fit(1e-2, 1)
```

Make way for augmentation by making the precomputations to False
`learn.precompute=False`

Fitting with Cycle length
`learn.fit(1e-2, 3, cycle_len=1)`

Unfreezing the layers
`learn.unfreeze()`

Differential Learning rates for layers
`lr=np.array([1e-4,1e-3,1e-2])`

Fitting with Cycle Multiplier
`learn.fit(lr, 3, cycle_len=1, cycle_mult=2)`

Storing and loading models
`learn.save('caavo_1')`
`learn.load('caavo_1')`

Making predictions on the test set
`log_preds = learn.predict(is_test=True)`
`log_preds = np.argmax(log_preds,axis=1)`
`print(log_preds)`

Getting the predictions for submission
`for i in range(log_preds.shape[0]):`
 `arr.append(data.classes[log_preds[i]])`
`print(arr)`

Submitting the solution csv file
`import pandas as pd`
`df2 = pd.DataFrame(arr)`
`df2.insert(0,'image_name',[item[5:] for item in data.test_ds.fnames])`
`df2.columns = ['image_name', 'category']`
`print(df2.head())`
`df2.to_csv('ANS1.csv',index=False)`


Submission ID:	Result	Score
16475611	✔ Accepted	0.45038
2 seconds ago		

22.	 veje	Veje Subash Gandyer	100.45 (1)
-----	---	----------------------------	---------------

Iteration 2 → Rank 20

Tried more epochs

Submission ID:	Result	Score
16476260	✓ Accepted	0.54327
3 seconds ago		

20.	 Veje Subash Gandyer veje	100.543 (1)	155:45:25	100 7:37:30	N/A
-----	---	----------------	-----------	----------------	-----

Iteration 3 → Rank 17

Tried more epochs, learning rates

Epoch	<div>100% 7/7 [13:48<00:00, 118.37s/it]</div>		
epoch	trn_loss	val_loss	accuracy
0	1.371284	1.972845	0.409412
1	1.363264	1.95151	0.415737
2	1.288256	1.947993	0.422247
3	1.276602	1.90089	0.425967
4	1.242705	1.934669	0.420108
5	1.179885	1.892373	0.437872
6	1.171088	1.926887	0.437686

3]: [1.9268874, 0.4376860111951828]

Submission ID:	Result	Score
16476745	✓ Accepted	0.56964
3 seconds ago		


Iteration 4 → Rank 15

More epochs, learning rates

Epoch 100% 7/7 [13:49<00:00, 118.46s/it]

epoch	trn_loss	val_loss	accuracy
0	1.172242	1.956882	0.422526
1	1.180492	1.913352	0.436756
2	1.119055	1.92889	0.446708
3	1.163991	1.932494	0.443917
4	1.095295	1.934552	0.446615
5	1.00317	1.994037	0.450707
6	0.963519	1.959744	0.453311

[1.9597443, 0.4533110111951828]

15.	 Vejey Subash Gandyer vejey	100.574 (1)	156:37:59	100 7:37:30	N/A
-----	--	----------------	-----------	----------------	-----


Iteration 5 → Rank 13

Epoch 100% 7/7 [13:50<00:00, 118.58s/it]

epoch	trn_loss	val_loss	accuracy
0	1.039096	1.962614	0.458054
1	1.033576	1.979503	0.46131
2	0.964347	2.002091	0.471075
3	1.043156	2.017644	0.466053
4	0.924182	2.032051	0.463914
5	0.88897	2.070672	0.465681
6	0.845514	2.047883	0.469773

.]: [2.0478833, 0.4697730665405591]

Submission ID: 16478212 8 seconds ago	Result ✔ Accepted	Score 0.58027
---	----------------------	------------------

13.	 Vejey Subash Gandyer vejey	100.58 (1)	157:08:27	100 7:37:30	N/A
-----	--	---------------	-----------	----------------	-----

Iteration 6 → Rank 5

Changing the size of image to 128 x 128 from 64 x 64

Epoch 100% 7/7 [14:23<00:00, 123.30s/it]

epoch	trn_loss	val_loss	accuracy
0	1.494176	1.88143	0.411551
1	1.386487	1.75352	0.456008
2	1.300985	1.757973	0.463821
3	1.259778	1.663669	0.484654
4	1.181627	1.663277	0.504836
5	1.142547	1.640895	0.521112
6	1.109125	1.635524	0.519159

: [1.6355238, 0.5191592251261076]

Submission ID:

16479262

3 seconds ago

Result

✓ Accepted

Score

0.63052

5.



Vejeey Subash Gandyer
vejeey

100.631
(1)

157:59:42

100
7:37:30

N/A

Iteration 7 → Rank 3

Epoch 100% 7/7 [14:22<00:00, 123.25s/it]

epoch	trn_loss	val_loss	accuracy
0	0.919342	1.578441	0.536086
1	0.957745	1.603818	0.534133
2	0.825132	1.609289	0.531343
3	0.901866	1.61382	0.538969
4	0.81104	1.620696	0.539342
5	0.769923	1.644314	0.548456
6	0.745644	1.639149	0.555618

: [1.6391493, 0.5556175584594408]

Submission ID:

16480199

3 seconds ago

Result

✓ Accepted

Score

0.65454

Leaderboard					
PROGRAMMERS		SCORE (200)	TOTAL TIME	P0	P1
View only in network		Problems Solved (2)	(HH:MM:SS)	Coin Game (100)	Caavo Comput... (100)
1.	 Biswesh Anupam Panda bisweshanupam	100.695 (1)	217:24:46	100 78:07:00	N/A
2.	 Akhil Kumar helloakhil	100.663 (1)	92:40:50	100 2:15:13	N/A
3.	 Veje Subash Gandyer veje	100.65 (1)	158:20:55	100 7:37:30	N/A
4.	 Suvronil Das suvronil	100.647 (1)	257:12:08	100 143:44:14	N/A
5.	 Gokula Krishnan GokulaKrishnan	100.632 (1)	101:09:47	100 48:01:11	N/A

That's the journey I made and reached 3rd place in the Computer Vision Challenge.

Edit: I ended up with 8th place at the end of the competition when I checked the leaderboard 2 days after the competition was over.